

Workflow Audit — Where AI Helps, Where It Doesn't

FlyRank AI-Fluency Track · FL-01 · Alex (Xiaoyu) Ji · 2026-07-12 · framework: Ethan Mollick, "Onboarding Your AI Intern"

JUST ME **DELEGATE + REVIEW** **COLLABORATE** **AUTOMATE (HITL)**

1 · The audit — 15 recurring tasks from my real week

Recurring task (real)	Classification	One-line rationale
Shipping features across my agent stack (Orallexa, Shadow, VibeX, Solo Founder OS)	Collaborate	I hold the architecture and judgment; AI accelerates implementation and edge-case coverage.
Pre-push code review + local build / typecheck	Delegate + review	AI flags diffs and regressions fast, but I approve before every push (build-quality-agent + my eyes).
Job-application cover letters & recruiter outreach	Collaborate	AI drafts to the JD; I supply the true claims and the final voice.
FlyRank internship assignments (ML + Fluency tracks)	Collaborate	The learning is mine; AI helps structure, run, and pressure-test the work.
Daily market / trading research briefs (Orallexa, Polymarket, SpaceX tracker)	Automate (HITL)	Cron-generated every day; I only read the mispricing flags — no per-run effort.
Bilingual EN / 中文 content sync & translation	Delegate + review	High-volume, pattern-heavy; I spot-check terminology and tone before it ships.
alex-brain memory vault & knowledge-index upkeep	Automate (HITL)	Indexing + source discovery run on cron; I curate weekly, not daily.
IEEE VR/VIS 2027 paper (co-first-author with Lora Levitchi)	Just me*	Co-authored, high-stakes, reputationally mine — AI may draft prose, but claims and framing stay under my hand.
OSS launch / marketing posts for my repos	Delegate + review	marketing-agent drafts platform-native copy; I verify facts (stars, test counts) and post.
Email replies (Lora, recruiters, professors)	Collaborate	AI drafts; I set the intent and never send without reading it myself.
Coursework & studying (MS CS, MCP & AI-agents course)	Just me	Understanding can't be delegated — AI can tutor me, but it can't learn it for me.
Debugging cron / infra failures (launchd, Vercel deploys)	Collaborate	AI hypothesizes fast, but I verify root cause before trusting a fix (verify-before-audit habit).
Real-money trading / sizing decisions	Just me	Irreversible, risk-bearing judgment — never automated; system stays paper-only until validated.
Resume tailoring for a target role	Collaborate	AI restructures to the role; I own the accuracy of every single line.
Recording demo Looms / screen walkthroughs	Just me	Founder voice and energy read as authentic; AI narration signals "non-founder" and gets discounted.

* IEEE paper is a hybrid: **just-me** ownership of claims/framing with **collaborate**-level drafting support. Counted as "just me" because the accountable decision is non-delegable. Clean "just me" tasks: coursework, real-money decisions, Loom recording.

2 · My three target tasks for FL-02 → FL-04 (with measurable "done well")

Target 1 — Tailored job-application cover letter / recruiter outreach.

Done well = aligned to ≥5 of the JD's named keywords; ≤1 page (outreach ≤150 words); every claim traceable to my resume/GitHub with no inflated numbers; ready to send after ≤2 human edits and under 15 minutes total.

Target 2 — Bilingual launch / marketing post for an OSS repo.

Done well = EN + 中文 versions at parity (same claims, platform-native phrasing); every factual number (stars, tests, endpoints) verified against the live repo before writing; publishable after ≤2 edits; zero fabricated features.

Target 3 — Turning a raw dataset/analysis into a framed, careful write-up (e.g., the FlyRank notebooks).

Done well = decision + action + cost-of-wrong-call named explicitly; ≥2 real numbers computed live from the data; careful language

throughout (observed / directional / decision-support, never "proven"); notebook runs top-to-bottom with 0 errors; reviewer-ready.

3 · Claude Project custom instructions (paste-ready — see appendix / .txt)

Configured a Claude Project named “**AI Fluency — Alex Ji**” with the instructions in the appendix below (who I am, tone preferences, current goals). Screenshot submitted separately.

4 · Evidence checklist (accounts + Academy)

- ☐ Claude account (have) · ☐ ChatGPT account · ☐ Anthropic Academy account
- ☐ Enrolled in *AI Fluency: Framework & Foundations* · ☐ Module 1 complete (screenshot)
- ☐ Claude Project created with custom instructions (screenshot)

Appendix · Claude Project custom instructions

Who I am: Alex (Xiaoyu) Ji — U.S. Navy veteran, MS Computer Science (AI specialization, GPA 4.0, Yeshiva University), AI/ML engineer. I build and ship production multi-agent LLM systems solo: Orallexa (AI trading OS), an 11-agent open-source "Solo Founder OS", Shadow (banking-compliance council), and VibexForge. Native bilingual English / 中文.

How I want you to work with me:

- Be direct and concise. Lead with the answer, then the reasoning. No filler, no flattery, no "great question."
- Reply in Chinese when I write in Chinese; keep code, technical terms, and any external-facing copy in English.
- Bias to action: when you have enough to act, act — don't ask me to confirm obvious defaults. Give a recommendation, not a menu.
- Verify claims against real sources (repos, data, files) before stating them; flag anything you couldn't verify. Never inflate numbers.
- Use careful language: observed / directional / decision-support — never "proven" or causal claims I can't back.
- Push back when I'm wrong.

Current goals (2026 H2):

- Land an AI/ML or forward-deployed engineer role (AI startups; veteran programs like Oracle OVIP).
- Complete the FlyRank internship (ML + AI-Fluency tracks) with portfolio-grade deliverables.
- Ship the IEEE VR/VIS 2027 co-first-author paper with Lora Levitchi (deadline 2026-08-24).
- Keep the open-source agent stack shipping and honestly benchmarked.