

Portfolio Cases — Alex (Xiaoyu) Ji

FlyRank · General AI Fluency · Week 2 "Frame It as Cases" · 2026-07-12

VOICE CARD: Direct, plain, honest — proof over adjectives, no buzzwords.

One audience: hiring managers / tech leads at AI startups · **One action:** book a 15-minute intro call

Work — case studies (one per sitemap piece)

Orallexa — an AI trading OS that reasons out loud

THE PROBLEM

Trading tools give you one of two bad things: a black-box "BUY" you can't audit, or a wall of indicators you have to interpret yourself. I wanted a system that argues its case out loud — and that physically can't do something reckless while I'm not looking.

WHAT I DID — AND DECIDED

I built Orallexa, where every call goes through a five-voice debate: Bull, Bear, and Judge argue each ticker, and a fourth agent — the Critic — audits the chain for same-source bias and missing base rates. Two decisions I stand behind: (1) the executor is *incapable* of placing a live trade — it's a code-level invariant asserted in a unit test, not a promise I make; (2) I route ~80% of calls to a cheap model and only escalate the hard reasoning, because paying frontier prices for boilerplate is how these systems quietly go broke.

WHAT CAME OF IT

54 stars on GitHub, 1,300+ passing tests, and per-analysis cost down 10× (\$0.05 → \$0.005) with no quality drop on a 47/50 blind review. It's the project I point to when someone asks "have you actually shipped an agent system?"

Solo Founder OS — the "AI employee" idea, but open and operator-controlled

THE PROBLEM

The AI-employee platforms (Devin, Lindy, Cofounder) are closed SaaS — you rent them, and they own your workflow. I wanted to see how much of a one-person company I could actually automate with open code I control.

WHAT I DID — AND DECIDED

I built an 11-agent open-source stack — one agent per real job of running a solo company: code review, marketing, customer discovery, support, payments, analytics. The decision that shaped it: agents can propose their own improvements, but a human approves every change through a GitHub PR. Self-improvement with a human gate — not a black box that rewrites itself overnight. Each agent is its own installable package, so you take one or all.

WHAT CAME OF IT

536 tests, published on PyPI, and it runs all eleven agents concurrently for under \$5/month. It's the closest thing I have to a thesis I can hand someone and say "clone it."

Shadow — a compliance council a bank can actually audit

THE PROBLEM

A bank can't ship an AI that says "approve this loan" and can't show its work — a regulator will ask why, and "the model decided" is not an answer.

WHAT I DID — AND DECIDED

I built Shadow with a banking-domain collaborator (Lora Levitchi, 14 years in banking software): a multi-voice council that returns a decision *plus* a tamper-evident audit chain, and exposes its tools over MCP so an analyst can query it from inside Cursor or Claude Desktop instead of curling an API. The decision I'm proudest of is one I *refused* to make — when I wired in the real risk thresholds my benchmark dropped a few points, and I published the drop instead of tuning the prompts to bury it. In a procurement review, that honesty is the moat.

WHAT CAME OF IT

A reproducible benchmark reported as 86 ± 1 (not one cherry-picked number), a hash-chain audit on every decision, and a FICO floor that's a hard block in code — not a suggestion.

About — bio

I'm Alex Ji — a U.S. Navy veteran turned AI engineer (MS Computer Science, GPA 4.0). I build production multi-agent systems solo and ship them open-source: a 54-star trading OS, an 11-agent stack for one-person companies, and a compliance council banks can actually audit. I care more about whether a system can show its work than whether the demo looks slick. I'm the engineer you hire to turn "we should use AI" into something running in production.

Contact — CTA

If you're building with AI and want someone who ships end-to-end — data, to agents, to the audit trail — let's talk. **Book a 15-minute call** [Calendly link] · or email xji1@mail.yu.edu.

Before / after — generic AI vs my voice

GENERIC AI LINE

"Results-driven AI engineer with a proven track record of leveraging cutting-edge machine learning to deliver impactful, scalable solutions across diverse domains."

MY EDITED VERSION

"I build multi-agent AI systems solo and ship them open-source. My trading OS has 54 stars and 1,300+ tests; my compliance council publishes its benchmark drops instead of hiding them. I'd rather show you the audit trail than the buzzwords."