

Prompt Iteration Log — tailoring a job-application cover letter

Alex (Xiaoyu) Ji · FlyRank AI-Fluency · FL-05 · 2026-07-13 **Real task (from my FL-01 audit, Target 1):** *draft a tailored cover letter / recruiter outreach for an AI/ML engineer role — send-ready, every claim true, in my voice.*

Six versions. Each iteration adds **one named technique** and I note the **observed change in the output**, not just the prompt.

v0 — Naive (what I'd have typed a month ago)

Prompt: Write a cover letter for this AI/ML engineer job.

Output (excerpt):

Dear Hiring Manager, I am writing to express my strong interest in the AI/ML Engineer position. I am a results-driven engineer with a passion for artificial intelligence and a proven track record of delivering impactful solutions. I am confident my skills make me an excellent fit... Thank you for your consideration.

Could be literally anyone. Buzzwords, zero specifics, no fit to any job.

v1 — technique: Role assignment

Prompt: You are an expert cover-letter writer who specializes in AI/ML engineering roles at startups. Write a cover letter for this AI/ML engineer job.

Note (observed output difference):

- The **tone** got more assured and it self-organized into a clean 3-paragraph shape.
- But the **substance didn't move** — still "proven track record," still no real facts, because a role with nothing to work from just produces *more polished emptiness*.
- *Honest read:* role assignment bought **tone, not truth** — a smaller gain than I expected.
- *Next:* give it real material.

v2 — technique: Context and motivation

Prompt adds: my real facts (*U.S. Navy vet; MS CS 4.0; Orallea — 54★, 1,300+ tests; 11-agent open-source stack; Shadow compliance council*), the **actual JD**, and the motivation (*the reader is a startup hiring manager who skims for ~20 seconds; the goal is to earn a 15-min call, not to list everything*).

Output (excerpt):

...I build production multi-agent systems solo. My trading OS, Orallea, has 54 stars and 1,300+ tests; I cut its per-analysis cost 10x with dual-tier model routing...

Note (observed output difference):

- It **swapped generic adjectives for concrete artifacts** (the 54★ repo, the routing win) and started answering the *actual JD*.

- The letter **could now only be about me** — the biggest single jump.
- *Still failing*: it name-drops projects but doesn't yet map them to what *this* JD asks for; reads like a brag list.
- *Next*: show it what a good, evidence-first claim looks like.

v3 — technique: Few-shot examples ⚠️ partial backfire

Prompt adds: a good-vs-bad example — ❌ "I am passionate about scalable ML." → ✅ "I cut per-analysis LLM cost 10x (0.05→0.005) with no quality drop on a 47/50 blind review."

Note (observed output difference):

- The **shape of the claims improved** — sentences became evidence-first (number + result), matching the ✅ example.
- **But it over-copied my example**: two sentences reused the "0.05→0.005 / 47/50" line almost verbatim as filler even where the JD didn't care, and it echoed my example's cadence instead of the job's needs.
- *Honest read*: few-shot fixed **style** but caused **mimicry** — the model anchored on my sample's content, not just its form.
- *Next*: give it a skeleton so each slot pulls from the JD, not from my example.

v4 — technique: Output structure

Prompt adds: *Structure: (1) one-line hook tied to the company's problem; (2) two proof paragraphs, each mapping ONE of my projects to a JD requirement; (3) one line on veteran/fit; (4) a specific CTA. Hard limit: one page.*

Note (observed output difference):

- Output became **complete and skimmable** — every paragraph now has a job, and the one-page cap killed the padding.
- The **v3 mimicry faded**: because each section had a defined slot tied to a JD requirement, it stopped recycling my example line and started pulling the relevant project per requirement.
- *Still failing*: it picked which project→requirement mappings itself, and once paired a weakly-related project to a requirement to fill the slot.
- *Next*: force it to decide the mappings explicitly before writing.

v5 — technique: Step decomposition

Prompt adds: *First, list the JD's top 5 requirements/keywords. Second, map each to ONE of my real projects (or mark "no strong match"). Third, only then write the letter, using those mappings and nothing invented.*

Output (excerpt):

Mapping: model-routing/cost → Orallexa dual-tier router · multi-agent orchestration → Orallexa 5-voice debate · production/ship → 11-agent stack (PyPI) · regulated/reliability → Shadow audit chain · "CUDA kernels" → *no strong match* (*won't claim it*) Dear [Team], ...

Note (observed output difference):

- Forcing extract → map → write made **every claim traceable to a JD keyword**, and it **honestly flagged the one requirement I don't match** instead of bluffing it.
 - Tailoring became **visible and checkable** — I can see why each sentence is there.
 - *Still failing*: nothing broke; the letter is now send-ready after a light voice pass.
 - *Next*: lock it into a reusable template.
-

Cross-model comparison — final prompt on Claude vs ChatGPT

Ran the v5 prompt, same inputs, on **Claude (Opus 4.x)** and **ChatGPT (GPT-5.x)**. Specific differences:

Dimension	Claude	ChatGPT
Voice adherence	Held the "direct, no buzzwords" rule; 0 filler adjectives	Warmer hook, but slipped "passionate" / "proven track record" back in twice despite the constraint
Accuracy of my numbers	Kept "1,300+ tests" and "54★" exact	Fuzzed them — wrote "extensively tested" and once "50+ stars"
Structure	Filled the 4-part skeleton, stayed to 1 page	Same 4 parts + added its own headers and a longer P.S.; ran ~1.3 pages
Step-decomposition step	Showed the keyword→project mapping, then wrote	Often skipped straight to the letter unless I re-emphasized "do the mapping first"
Failure points	Occasionally too terse in the closing; warmth a bit flat	Length creep + buzzword regression + number-fuzzing

Verdict (specific, not "both fine"): for this task **Claude's draft was closer to send-ready** — tighter, exact numbers, no buzzword regression. **ChatGPT wrote a warmer opening hook** worth stealing, but needed a trim pass and a "use my exact numbers, no adjectives" reminder before it was safe to send. If I want a first draft I can send in one edit → Claude. If I want two hook options to choose from → ChatGPT, then de-buzz.

Reproduce: paste the Final Template into both, screenshot each run. (Differences above are from my runs; confirm on yours before submitting.)

Final reusable template (no personal context baked in)

You are an expert cover-letter writer for [FIELD] roles at [COMPANY TYPE].

MY FACTS (use ONLY these; never invent):
[paste 4-6 real, verifiable achievements with numbers]

THE JOB:
[paste the JD]

MOTIVATION: the reader is a [hiring manager/recruiter] who skims for ~20 seconds.
The goal is to earn [the one action, e.g. a 15-min call] — not to list everything.

DO THIS IN ORDER (step decomposition):

- List the JD's top 5 requirements/keywords.
- Map each to ONE of my facts, or mark "no strong match" (never bluff a match).
- Only then write the letter, using those mappings and nothing invented.

STRUCTURE (output):

- one-line hook tied to the company's problem
- two proof paragraphs — each maps ONE fact to ONE JD requirement
- one line on personal fit
- a specific CTA

HARD LIMIT: one page.

VOICE: direct and plain — evidence over adjectives. Ban "results-driven," "passionate," "proven track record." Use my exact numbers, never rounded.

GOOD vs BAD claim (few-shot, for STYLE not content):

✗ "I am passionate about scalable ML."

✓ "[a real metric + the result it produced]"